



YouTube Shorts Performance Prediction

Supervised Machine Learning Case Study

Objective: Predict engagement performance (Low / Medium / High) of YouTube Shorts using content and publishing features.

```
In [ ]: !pip install catboost
```

Collecting catboost

Downloading catboost-1.2.8-cp312-cp312-manylinux2014_x86_64.whl.metadata (1.2 kB)
Requirement already satisfied: graphviz in /usr/local/lib/python3.12/dist-packages (from catboost) (0.21)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from catboost) (3.10.0)
Requirement already satisfied: numpy<3.0,>=1.16.0 in /usr/local/lib/python3.12/dist-packages (from catboost) (2.0.2)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.12/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from catboost) (1.16.3)
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.12/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.24->catboost) (2025.3)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (1.3.3)
Requirement already satisfied: cyclers>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (4.61.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->catboost) (3.3.1)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly->catboost) (9.1.2)
Downloading catboost-1.2.8-cp312-cp312-manylinux2014_x86_64.whl (99.2 MB)

99.2/99.2 MB 8.7 MB/s eta 0:00:00
Installing collected packages: catboost
Successfully installed catboost-1.2.8

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import gdown
import re

from typing import Dict, List, Sequence, Optional

from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier

from sklearn.preprocessing import LabelEncoder
from scipy.stats import randint, uniform

import warnings
warnings.filterwarnings('ignore')
sns.set(style='whitegrid')
```

Data Loading , Cleaning & EDA

Problem Statement & Objective:

From a business perspective, accurately predicting High-engagement Shorts enables creators to optimize video length, title structure, category selection, and upload timing before publishing, thereby improving reach, growth, and monetization potential.

Load Dataset (Upload CSV)

```
In [ ]: !gdown --id '1oSKX150eAXFDLtw34gfUg0hRDBkYZE5s'
```

/usr/local/lib/python3.12/dist-packages/gdown/__main__.py:140: FutureWarning: Option `--id` was deprecated in version 4.3.1 and will be removed in 5.0. You don't need to pass it anymore to use a file ID.

warnings.warn(
Downloading...

From: https://drive.google.com/uc?id=1oSKX150eAXfDLtW34gfUgOhRDBkYZE5s

To: /content/youtube_shorts_performance_dataset.csv

100% 18.7k/18.7k [00:00<00:00, 31.8MB/s]

```
In [ ]: df = pd.read_csv('/content/youtube_shorts_performance_dataset.csv')  
df.head()
```

```
Out[ ]:   video_id  title  duration_sec  hashtags_count  views  likes  comments  shares  upload
```

	video_id	title	duration_sec	hashtags_count	views	likes	comments	shares	upload
0	vid_1000	Short Video #0	43	9	198775	21933	3228	400	
1	vid_1001	Short Video #1	56	2	290336	20063	3719	1942	
2	vid_1002	Short Video #2	33	6	264206	37032	3228	1817	
3	vid_1003	Short Video #3	19	9	85076	27269	2371	980	
4	vid_1004	Short Video #4	47	8	90780	8041	2891	1109	

Data Cleaning and EDA overview

```
In [ ]: def eda_overview(df: pd.DataFrame) -> None:  
    print("
```

📐 Shape: (300, 10)

📄 Info

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300 entries, 0 to 299
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   video_id            300 non-null   object
1   title               300 non-null   object
2   duration_sec        300 non-null   int64
3   hashtags_count      300 non-null   int64
4   views               300 non-null   int64
5   likes               300 non-null   int64
6   comments            300 non-null   int64
7   shares              300 non-null   int64
8   upload_hour         300 non-null   int64
9   category            300 non-null   object
dtypes: int64(7), object(3)
memory usage: 23.6+ KB
```

❓ Missing Values

missing_count	
video_id	0
title	0
duration_sec	0
hashtags_count	0
views	0
likes	0
comments	0
shares	0
upload_hour	0
category	0

📊 Describe (numeric)

	duration_sec	hashtags_count	views	likes	comments	shares
count	300.000000	300.000000	300.000000	300.000000	300.000000	300.000000
mean	31.730000	4.343333	245058.166667	23063.126667	2638.296667	999.576667
std	16.042912	2.956562	141338.982093	14597.458220	1416.827287	564.410105
min	5.000000	0.000000	1404.000000	109.000000	40.000000	2.000000
25%	18.000000	2.000000	129620.500000	10342.000000	1445.750000	516.250000
50%	32.000000	4.000000	255962.000000	21779.500000	2863.500000	988.000000
75%	45.000000	7.000000	356805.000000	36706.500000	3761.000000	1463.750000
max	59.000000	9.000000	499401.000000	49923.000000	4971.000000	1998.000000

 Describe (categorical)

	video_id	title	category
count	300	300	300
unique	300	300	6
top	vid_1299	Short Video #299	Food
freq	1	1	57

Insight :

This model enables creators to optimize video length, titles, and upload timing before publishing, increasing the probability of achieving High engagement and therefore higher reach, subscriber growth, and monetization opportunities

Standardize Columns(snake_case)

```
In [ ]: def is_snake_case(col_name: str) -> bool:
        """
        Returns True if a column name is in snake_case.
        """
        pattern = r'^[a-z][a-z0-9]*[a-z0-9]$\n        return bool(re.match(pattern, col_name))
```

```
In [ ]: def dataframe_has_snake_case_columns(df):
        """
        Checks whether all DataFrame columns are in snake_case.
        """
        return all(is_snake_case(col) for col in df.columns)
```

```
In [ ]: def standardize_columns_to_snake_case(df):
        """
        Standardizes pandas DataFrame column names to snake_case.
```

```

"""
df = df.copy()

df.columns = (
    df.columns
    .str.strip()
    .str.lower()
    .str.replace(' ', '_')
    .str.replace('.', '_', regex=False)
    .str.replace('-', '_', regex=False)
)

return df

```

```

In [ ]: def ensure_snake_case(df):
        """
        Ensures DataFrame columns are snake_case.
        Applies transformation only if needed.
        """

        if dataframe_has_snake_case_columns(df):
            print("✅ Columns already in snake_case")
            return df
        else:
            print("🔪 Converting columns to snake_case")
            print("Standardized Columns:", df.columns.tolist())
            return standardize_columns_to_snake_case(df)

```

```

In [ ]: print("🔍 Checking column names for snake_case...")
        df = ensure_snake_case(df)

```

```

🔍 Checking column names for snake_case...
✅ Columns already in snake_case

```

Missing values

```

In [ ]: def fill_missing_values(df: pd.DataFrame, verbose: bool = True) -> pd.DataFrame:
        """
        Fills missing values:
        - Numeric columns -> median
        - Categorical columns -> mode
        Applies only if missing values exist.
        """

        df = df.copy()
        missing_total = df.isnull().sum().sum()

        if missing_total == 0:
            if verbose:
                print("✅ No missing values found. Nothing to fill.")
            return df

        if verbose:
            print(f"⚠️ Missing values detected: {missing_total}")
            print("🔧 Applying imputation...")

        # Numeric columns

```

```

num_cols = df.select_dtypes(include="number").columns
for col in num_cols:
    if df[col].isnull().any():
        df[col] = df[col].fillna(df[col].median())

# Categorical columns
cat_cols = df.select_dtypes(include="object").columns
for col in cat_cols:
    if df[col].isnull().any():
        df[col] = df[col].fillna(df[col].mode()[0])

if verbose:
    print("✅ Missing values filled successfully.")

return df

```

```
In [ ]: df = fill_missing_values(df)
```

✅ No missing values found. Nothing to fill.

Target Variable Creation

```

In [ ]: def add_engagement_metrics(
    df: pd.DataFrame,
    likes_col: str = "likes",
    comments_col: str = "comments",
    shares_col: str = "shares",
    views_col: str = "views",
) -> pd.DataFrame:
    """
    Adds engagement_rate and performance_engagement_tertile columns.

    Tertiles are based on 33rd and 66th percentiles.
    Returns a new DataFrame (does not mutate input).
    """

    required_cols = {likes_col, comments_col, shares_col, views_col}
    missing_cols = required_cols - set(df.columns)

    if missing_cols:
        raise ValueError(f"Missing required columns: {missing_cols}")

    df_out = df.copy()

    # ---- Safe engagement rate calculation ----
    df_out["engagement_rate"] = np.where(
        df_out[views_col] > 0,
        (
            df_out[likes_col]
            + df_out[comments_col]
            + df_out[shares_col]
        ) / df_out[views_col],
        np.nan, # avoid division by zero
    )

```

```

# ---- Compute tertile thresholds (ignore NaNs) ----
q33, q66 = df_out["engagement_rate"].quantile([0.33, 0.66])

# ---- Assign tertiles ----
df_out["performance_engagement_tertile"] = pd.cut(
    df_out["engagement_rate"],
    bins=[-np.inf, q33, q66, np.inf],
    labels=["Low", "Medium", "High"],
    include_lowest=True,
)
print("Target Variable Creation successfully.")
return df_out

```

```
In [ ]: df = add_engagement_metrics(df)
```

Target Variable Creation successfully.

Target distribution

```

In [ ]: def display_target_distribution(
    df: pd.DataFrame,
    target_col: str,
    plot: bool = True
) -> pd.DataFrame:
    """
    Displays count and proportion of target variable.
    Optionally plots a bar chart.
    """

    dist = (
        df[target_col]
        .value_counts(dropna=False)
        .to_frame("count")
    )

    dist["proportion"] = (dist["count"] / dist["count"].sum()).round(3)

    print( Target Distribution")
    #display(dist)


    if plot:
        plt.figure()
        plt.bar(dist.index.astype(str), dist["proportion"])
        plt.xlabel(target_col)
        plt.ylabel("Proportion")
        plt.title(f"Distribution of {target_col}")
        plt.ylim(0, 1)
        plt.show()

    return dist

```

```
In [ ]: display_target_distribution(  
    df,  
    target_col="performance_engagement_tertile"  
)
```

Target Distribution



```
Out[ ]:
```

	count	proportion
performance_engagement_tertile		
High	102	0.34
Low	99	0.33
Medium	99	0.33

Key insights :

The tertile-based target variable results in a near-balanced class distribution across Low, Medium, and High engagement categories, minimizing severe class imbalance issues. However, overlap between Medium and adjacent classes necessitates the use of F1-macro as the primary evaluation metric.

🚖 Imbalance check

```
In [ ]: def check_class_imbalance(
    series: pd.Series,
    severe_threshold: float = 0.10,
    moderate_threshold: float = 0.20,
    verbose: bool = True
) -> dict:
    """
    Checks class imbalance severity and prints result if verbose=True.
    """

    proportions = series.value_counts(normalize=True)

    min_ratio = proportions.min()
    max_ratio = proportions.max()
    imbalance_ratio = max_ratio / min_ratio

    if min_ratio < severe_threshold:
        severity = "Severe Imbalance"
        emoji_msg = "🔴 Severe imbalance detected"
    elif min_ratio < moderate_threshold:
        severity = "Moderate Imbalance"
        emoji_msg = "🟡 Moderate imbalance detected"
    else:
        severity = "Balanced"
        emoji_msg = "✅ Data is balanced"

    report = {
        "severity": severity,
        "min_class_ratio": round(min_ratio, 3),
        "max_class_ratio": round(max_ratio, 3),
        "imbalance_ratio": round(imbalance_ratio, 2),
        "class_proportions": proportions.to_dict(),
    }

    if verbose:
        print(emoji_msg)
        print(f"Min class ratio : {report['min_class_ratio']}")
        print(f"Max class ratio : {report['max_class_ratio']}")
        print(f"Imbalance ratio : {report['imbalance_ratio']}")

    return report
```

```
In [ ]: imbalance_report = check_class_imbalance(
    df["performance_engagement_tertile"]
)
```

```
✅ Data is balanced
Min class ratio : 0.33
Max class ratio : 0.34
Imbalance ratio : 1.03
```

Insight :

The tertile split produces a roughly balanced target (~33% per class), reducing severe class imbalance. However, Medium often overlaps with Low/High, justifying the use of F1-macro

rather than accuracy.



Feature Engineering & Transformation

```
In [ ]: def compute_peak_hours(
    df: pd.DataFrame,
    hour_col: str = "upload_hour",
    engagement_col: str = "engagement_rate",
    top_percentile: float = 0.75,
    min_samples_per_hour: int = 10,
    verbose: bool = True,
    plot: bool = False
) -> Dict[str, object]:
    """
    Identifies peak hours based on mean engagement rate.
    Optionally visualizes hourly engagement.

    Returns
    -----
    dict with keys:
        - peak_hours: List[int]
        - hourly_stats: pd.DataFrame
        - threshold: float
    """

    # -----
    # Validation
    # -----
    if not 0 < top_percentile < 1:
        raise ValueError("top_percentile must be between 0 and 1")

    required_cols = {hour_col, engagement_col}
    missing_cols = required_cols - set(df.columns)
    if missing_cols:
        raise ValueError(f"Missing required columns: {missing_cols}")

    # -----
    # Prepare valid data
    # -----
    df_valid = (
        df[[hour_col, engagement_col]]
        .dropna()
        .copy()
    )

    df_valid = df_valid[df_valid[hour_col].between(0, 23)]

    if df_valid.empty:
        raise ValueError("No valid rows after cleaning hour/engagement data")

    # -----
    # Hourly aggregation
    # -----
    hourly_stats = (
```

```

df_valid
.groupby(hour_col)
.agg(
    mean_engagement=(engagement_col, "mean"),
    samples=(engagement_col, "count")
)
.sort_index()
)

# -----
# Filter low-sample hours
# -----

hourly_stats = hourly_stats[
    hourly_stats["samples"] >= min_samples_per_hour
]

if hourly_stats.empty:
    raise ValueError(
        f"No hours meet min_samples_per_hour={min_samples_per_hour}"
    )

# -----
# Compute threshold and peak hours
# -----

threshold = hourly_stats["mean_engagement"].quantile(top_percentile)

peak_hours: List[int] = (
    hourly_stats
    .loc[hourly_stats["mean_engagement"] >= threshold]
    .index
    .tolist()
)

# -----
# Optional logging
# -----

if verbose:
    print("📊 Peak Hour Analysis")
    print(f"Top percentile      : {int(top_percentile * 100)}%")
    print(f"Engagement threshold : {round(threshold, 6)}")
    print(f"Peak hours           : {peak_hours}")

# -----
# Optional visualization (EDA use)
# -----

if plot:
    plt.figure()
    plt.plot(
        hourly_stats.index,
        hourly_stats["mean_engagement"],
        marker="o"
    )
    plt.axhline(
        threshold,
        linestyle="--",
        label="Peak Threshold"
    )

```

```

    )
    plt.xlabel("Upload Hour")
    plt.ylabel("Average Engagement Rate")
    plt.title("📊 Hourly Engagement Pattern")
    plt.xticks(range(0, 24))
    plt.legend()
    plt.grid(True)
    plt.show()

    return {
        "peak_hours": peak_hours,
        "hourly_stats": hourly_stats,
        "threshold": threshold
    }

```

```

In [ ]: # 1. Compute peak hours
peak_info = compute_peak_hours(df, plot=True)

# 2. Extract peak hours
peak_hours = peak_info["peak_hours"]

hourly_stats = peak_info["hourly_stats"]


```

📊 Peak Hour Analysis

Top percentile : 75%

Engagement threshold : 0.43106

Peak hours : [3, 5, 10, 11, 21]



```

In [ ]: def add_derived_features(
        df: pd.DataFrame,
        peak_hours: Optional[Sequence[int]] = None,
        title_col: str = "title",
        duration_col: str = "duration_sec",
        upload_hour_col: str = "upload_hour",
        log_cols: Sequence[str] = ("views", "likes", "comments", "shares"),
        verbose: bool = True
    ) -> pd.DataFrame:
        """
        Adds derived features:
        - Textual features from title
        - Rate-based features (per second)
        - Log-transformed numeric features
        - Time-based peak-hour indicator

        Returns a new DataFrame (does not mutate input).
        """

        # -----
        # Validation
        # -----
        required_cols = {title_col, duration_col, upload_hour_col, *log_cols}
        missing_cols = required_cols - set(df.columns)

        if missing_cols:
            raise ValueError(f"Missing required columns: {missing_cols}")

        if verbose:
            print("✳ Adding derived features...")
            print(f"Input rows : {df.shape[0]}")
            print(f"Input cols : {df.shape[1]}")

        df_out = df.copy()

        # -----
        # 1. Textual Features
        # -----
        title_series = df_out[title_col].fillna("")

        df_out["title_len_chars"] = title_series.str.len()
        df_out["title_word_count"] = title_series.str.split().str.len()
        df_out["title_has_question_mark"] = (
            title_series.str.contains(r"\?", regex=True).astype("int8")
        )

        if verbose:
            print("📄 Text features added: ['title_len_chars', 'title_word_count', 'title_has_question_mark']")

        # -----
        # 2. Rate Features
        # -----
        positive_duration = df_out[duration_col] > 0

        for metric in ("likes", "comments", "shares"):

```

```

        df_out[f"{metric}_per_sec"] = np.where(
            positive_duration,
            df_out[metric] / df_out[duration_col],
            np.nan
        )

    if verbose:
        print("🕒 Rate features added: ['likes_per_sec', 'comments_per_sec', 'shares_per_sec']")

    # -----
    # 3. Logarithmic Transformations
    # -----
    for col in log_cols:
        df_out[f"log_{col}"] = np.log1p(
            df_out[col].clip(lower=0)
        )

    if verbose:
        print(f"📈 Log features added: {list(log_cols)}")

    # -----
    # 4. Time-Based Feature (Peak Hour)
    # -----
    if peak_hours is not None:
        df_out["is_peak_hour"] = (
            df_out[upload_hour_col]
            .isin(peak_hours)
            .astype("int8")
        )

        if verbose:
            print(f"🕒 Peak-hour feature added (hours={list(peak_hours)})")
    else:
        df_out["is_peak_hour"] = 0

        if verbose:
            print("🕒 Peak-hour feature skipped (no peak hours provided)")

    if verbose:
        print("✅ Derived feature generation complete")
        print(f"📊 Output cols : {df_out.shape[1]}")

    return df_out

```

In []: df = add_derived_features(df, peak_hours)

✚ Adding derived features...

Input rows : 300

Input cols : 23

📄 Text features added: ['title_len_chars', 'title_word_count', 'title_has_question_mark']

🕒 Rate features added: ['likes_per_sec', 'comments_per_sec', 'shares_per_sec']

📈 Log features added: ['views', 'likes', 'comments', 'shares']

🕒 Peak-hour feature added (hours=[3, 5, 10, 11, 21])

✅ Derived feature generation complete

Output cols : 23

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300 entries, 0 to 299
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   video_id                             300 non-null    object
1   title                                300 non-null    object
2   duration_sec                         300 non-null    int64
3   hashtags_count                      300 non-null    int64
4   views                               300 non-null    int64
5   likes                               300 non-null    int64
6   comments                            300 non-null    int64
7   shares                              300 non-null    int64
8   upload_hour                         300 non-null    int64
9   category                             300 non-null    object
10  engagement_rate                     300 non-null    float64
11  performance_engagement_tertile      300 non-null    category
12  title_len_chars                     300 non-null    int64
13  title_word_count                    300 non-null    int64
14  title_has_question_mark             300 non-null    int8
15  likes_per_sec                       300 non-null    float64
16  comments_per_sec                    300 non-null    float64
17  shares_per_sec                      300 non-null    float64
18  log_views                           300 non-null    float64
19  log_likes                           300 non-null    float64
20  log_comments                        300 non-null    float64
21  log_shares                          300 non-null    float64
22  is_peak_hour                        300 non-null    int8
dtypes: category(1), float64(8), int64(9), int8(2), object(3)
memory usage: 48.0+ KB
```

```
In [ ]: def assert_derived_features_created(
    df: pd.DataFrame,
    expected_features: Sequence[str],
    verbose: bool = True
) -> None:
    """
    Asserts that all expected derived features exist in the DataFrame.
    Raises AssertionError if any feature is missing.
    """

    missing_features = set(expected_features) - set(df.columns)

    if missing_features:
        raise AssertionError(
            f"❌ Missing derived features: {sorted(missing_features)}"
        )

    if verbose:
        print(f"✅ All derived features created successfully")
        print(f"Total validated features : {len(expected_features)}")
```

```
In [ ]: EXPECTED_DERIVED_FEATURES = [
    "title_len_chars",
    "title_word_count",
    "title_has_question_mark",
    "likes_per_sec",
    "comments_per_sec",
    "shares_per_sec",
    "log_views",
    "log_likes",
    "log_comments",
    "log_shares",
    "is_peak_hour"
]
```

```
In [ ]: assert_derived_features_created(
    df,
    expected_features=EXPECTED_DERIVED_FEATURES
)
```

✓ All derived features created successfully
Total validated features : 11

```
In [ ]: #validate all record is_peak_hour column is binary
assert set(df["is_peak_hour"].unique()) <= {0, 1}
```

Feature engineering was guided by viewer behavior hypotheses:

- Title length and word count capture curiosity and click intent
- Question marks act as engagement triggers
- Rate-based features normalize engagement for video length
- Log transformations reduce skewness and outlier impact



Exploratory Data Analysis (EDA) & Feature Insights

The objective of EDA is to understand engagement behavior patterns, identify key drivers of High performance, and inform feature engineering and model selection decisions.

Common function for get all plots

```
In [ ]: # -----
# Global plot style (consistent across project)
# -----
def set_plot_style():
    sns.set_theme(style="whitegrid")
    plt.rcParams.update({
        "figure.dpi": 120,
        "axes.titlesize": 12,
        "axes.labelsize": 10,
        "xtick.labelsize": 9,
```

```

        "ytick.labelsize": 9,
    })

# -----
# 1. Title Length vs Engagement Rate
# -----
def plot_title_length_vs_engagement(
    df: pd.DataFrame,
    save_path: Optional[str] = None
):
    required_cols = {"title_len_chars", "engagement_rate"}
    if not required_cols.issubset(df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    plt.figure(figsize=(6, 4))
    sns.regplot(
        data=df,
        x="title_len_chars",
        y="engagement_rate",
        scatter_kws={"alpha": 0.4},
        line_kws={"color": "red"},
    )

    plt.title("Title Length vs Engagement Rate")
    plt.xlabel("Title Length (Characters)")
    plt.ylabel("Engagement Rate")

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

# -----
# 2. Engagement Rate Distribution Histogram
# -----

#Histogram + Tertile Cutoffs + Class Counts
#-----

def plot_engagement_distribution(
    df: pd.DataFrame,
    bins: int = 30,
    show_tertiles: bool = True,
    save_path: Optional[str] = None
):
    if "engagement_rate" not in df.columns:
        raise ValueError("Missing column: engagement_rate")

    plt.figure(figsize=(6, 4))
    data = df["engagement_rate"].dropna()

    sns.histplot(
        data.dropna(),
        bins=bins,
        kde=True,
    )

```

```

    if show_tertiles:
        q33, q66 = data.quantile([0.33, 0.66])
        plt.axvline(q33, color="orange", linestyle="--", label="33rd Percentile")
        plt.axvline(q66, color="red", linestyle="--", label="66th Percentile")
        plt.legend()

    plt.title("Engagement Rate Distribution")
    plt.xlabel("Engagement Rate")
    plt.ylabel("Frequency")

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

# -----
# 3. Boxplot: Engagement Rate vs Category
# -----
def plot_engagement_by_category(
    df: pd.DataFrame,
    save_path: Optional[str] = None
):
    required_cols = {
        "category",
        "engagement_rate",
        "performance_engagement_tertile"
    }
    if not required_cols.issubset(df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    plt.figure(figsize=(8, 4))
    sns.boxplot(
        data=df,
        x="category",
        y="engagement_rate",
        hue="performance_engagement_tertile",
    )

    plt.title("Engagement Rate by Category & Performance")
    plt.xlabel("Category")
    plt.ylabel("Engagement Rate")
    plt.xticks(rotation=45)
    plt.legend(title="Performance", loc="best")

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

# -----
# 4. Correlation Heatmap
# -----
def plot_correlation_heatmap(
    df: pd.DataFrame,
    save_path: Optional[str] = None
):

```

```

numeric_df = df.select_dtypes(include="number")

if numeric_df.empty:
    raise ValueError("No numeric columns found for correlation")

corr_matrix = numeric_df.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(
    corr_matrix,
    cmap="coolwarm",
    center=0,
    linewidths=0.5,
    cbar_kws={"shrink": 0.8},
)

plt.title("Correlation Heatmap of Numerical Features")

if save_path:
    plt.savefig(save_path, bbox_inches="tight")
plt.show()

# -----
# 5. Upload Hour vs Average Engagement Rate
# -----
def plot_hourly_engagement(
    df: pd.DataFrame,
    save_path: Optional[str] = None
):
    required_cols = {"upload_hour", "engagement_rate"}
    if not required_cols.issubset(df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    hourly_engagement = (
        df
        .dropna(subset=["engagement_rate"])
        .groupby("upload_hour")["engagement_rate"]
        .mean()
        .sort_index()
    )

    plt.figure(figsize=(10, 4))

    colors = [
        "orange" if h in peak_hours else "steelblue"
        for h in hourly_engagement.index
    ]

    hourly_engagement.plot(kind="bar", color=colors)

    plt.title("Average Engagement Rate by Upload Hour\n(Peak Hours Highlighted)")
    plt.xlabel("Upload Hour (0-23)")
    plt.ylabel("Average Engagement Rate")

```

```

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

# -----
# 6. Scatter Plot: Duration vs Engagement Rate
# -----
def plot_duration_vs_engagement(
    df: pd.DataFrame,
    save_path: Optional[str] = None
):
    required_cols = {
        "duration_sec",
        "engagement_rate",
        "performance_engagement_tertile"
    }
    if not required_cols.issubset(df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    plt.figure(figsize=(6, 4))
    sns.scatterplot(
        data=df,
        x="duration_sec",
        y="engagement_rate",
        hue="performance_engagement_tertile",
        alpha=0.6,
    )

    plt.title("Duration vs Engagement Rate")
    plt.xlabel("Duration (seconds)")
    plt.ylabel("Engagement Rate")

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

```

In []: peak_hours

Out[]: [3, 5, 10, 11, 21]

Print Eda Header

```

In [ ]: def print_eda_header(
    section_title: str,
    section_number: int | None = None
) -> None:
    """
    Prints a formatted EDA section header.

    Example:
    ### 1. Title Length vs Engagement Rate
    """

```

```

if section_number is not None:
    print(f"### {section_number}. {section_title}")
else:
    print(f"### {section_title}")

print() # spacing

```

Print Eda Section

```

In [ ]: def print_eda_section(
        insights: list[str],
        implications: list[str]
    ) -> None:
    """
    Prints formatted EDA insights and implications.
    """

    if insights:
        print("🔍 Insight:")
        for insight in insights:
            print(f"- {insight}")
        print()

    if implications:
        print("🔗 Implication:")
        for implication in implications:
            print(f"- {implication}")
        print()

    print("-" * 70)

```

Generate Eda Report

```

In [ ]: def GenerateEdaReport(df):
    print("📊 Exploratory Data Analysis (EDA) Report")

    # -----
    # Apply global plot style
    # -----
    set_plot_style()

    # =====
    # 1. Title Length vs Engagement Rate
    # =====
    print_eda_header(
        section_number=1,
        section_title="Title Length vs Engagement Rate"
    )

    plot_title_length_vs_engagement(df)

    print_eda_section(
        insights=[

```

```

        "Engagement does not increase linearly with title length.",
        "Most titles cluster between 14-16 characters with low-to-moderate enga
        "High-engagement outliers suggest title length alone is not a strong dr
        "Titles with moderate length and questions increase curiosity, improvin
    ],
    implications=[
        "Title length should be treated as a supporting feature.",
        "Its impact likely depends on interaction with other content attributes
    ]
)

# =====
# 2. Engagement Rate Distribution
# =====
print_eda_header(
    section_number=2,
    section_title="Engagement Rate Distribution"
)

plot_engagement_distribution(df)

print_eda_section(
    insights=[
        "The engagement rate distribution is highly right-skewed.",
        "Most content receives very low engagement, with a long tail of high pe
    ],
    implications=[
        "Log transformation is justified to reduce skewness.",
        "Binning engagement into tertiles provides a more robust target."
    ]
)

# =====
# 3. Engagement Rate vs Category
# =====
print_eda_header(
    section_number=3,
    section_title="Engagement Rate vs Category"
)

plot_engagement_by_category(df)

print_eda_section(
    insights=[
        "Engagement varies significantly across content categories.",
        "Lifestyle, Travel, and Education show higher median engagement.",
        "Clear separation across performance tertiles validates target labeling
        "Comedy and Lifestyle exhibit both high median engagement and lower va
    ],
    implications=[
        "Category is a strong predictive feature.",
        "Category-specific behavior should be captured in feature engineering."
    ]
)

# =====

```

```

# 4. Correlation Analysis
# =====
print_eda_header(
    section_number=4,
    section_title="Correlation Analysis"
)

plot_correlation_heatmap(df)

print_eda_section(
    insights=[
        "Raw engagement metrics are strongly correlated, indicating multicollin",
        "Normalized and log-scaled features show more stable relationships.",
        "Duration and hashtag count show weak standalone correlation.",
        "Hashtags beyond a moderate count show diminishing returns, indicating
    ],
    implications=[
        "Redundant engagement signals should be handled carefully.",
        "Normalized or transformed metrics are preferable for modeling."
    ]
)

# =====
# 5. Upload Hour vs Engagement Rate
# =====
print_eda_header(
    section_number=5,
    section_title="Upload Hour vs Engagement Rate"
)

plot_hourly_engagement(df)

print_eda_section(
    insights=[
        "Engagement varies notably by upload hour.",
        "Early morning and late evening hours show engagement peaks.",
        "Shorts uploaded between 18-21 show ~20-30% higher average engagement t
    ],
    implications=[
        "Posting time has a measurable impact on engagement.",
        "Peak-hour indicators are justified as contextual features."
    ]
)

# =====
# 6. Duration vs Engagement Rate
# =====
print_eda_header(
    section_number=6,
    section_title="Duration vs Engagement Rate"
)

plot_duration_vs_engagement(df)

print_eda_section(
    insights=[

```

```


    "There is no strong linear relationship between duration and engagement",
    "High-engagement outliers appear at mid-range durations.",
    "Very long videos do not consistently outperform shorter ones.",
    "Shorts within the optimal duration range maximize completion rates,whi
],
implications=[
    "Duration should be treated as a supporting, non-linear feature.",
    "Its effect likely depends on content type and context."
]
)

```

GenerateCompleteEDA Report

In []: `GenerateEdaReport(df)`

Exploratory Data Analysis (EDA) Report
 ### 1. Title Length vs Engagement Rate




🔍 Insight:

- Engagement does not increase linearly with title length.
- Most titles cluster between 14-16 characters with low-to-moderate engagement.
- High-engagement outliers suggest title length alone is not a strong driver.
- Titles with moderate length and questions increase curiosity, improving click-through and early engagement signals that boost Shorts distribution

🎯 Implication:

- Title length should be treated as a supporting feature.
- Its impact likely depends on interaction with other content attributes.

2. Engagement Rate Distribution




🔍 Insight:

- The engagement rate distribution is highly right-skewed.
- Most content receives very low engagement, with a long tail of high performers.

🎯 Implication:


- Log transformation is justified to reduce skewness.
- Binning engagement into tertiles provides a more robust target.

3. Engagement Rate vs Category



-  Insight:
- Engagement varies significantly across content categories.
 - Lifestyle, Travel, and Education show higher median engagement.
 - Clear separation across performance tertiles validates target labeling.
 - Comedy and Lifestyle exhibit both high median engagement and lower variance, making them more reliable categories for sustained performance.
-  Implication:
- Category is a strong predictive feature.
 - Category-specific behavior should be captured in feature engineering.

4. Correlation Analysis




🔍 Insight:

- Raw engagement metrics are strongly correlated, indicating multicollinearity.
- Normalized and log-scaled features show more stable relationships.
- Duration and hashtag count show weak standalone correlation.
- Hashtags beyond a moderate count show diminishing returns, indicating that excessive hashtag usage does not materially improve engagement.

🔗 Implication:

- Redundant engagement signals should be handled carefully.
- Normalized or transformed metrics are preferable for modeling.

5. Upload Hour vs Engagement Rate



 Insight:

- Engagement varies notably by upload hour.
- Early morning and late evening hours show engagement peaks.
- Shorts uploaded between 18-21 show ~20-30% higher average engagement than morning uploads.

 Implication:

- Posting time has a measurable impact on engagement.
- Peak-hour indicators are justified as contextual features.

6. Duration vs Engagement Rate



🔍 Insight:

- There is no strong linear relationship between duration and engagement.
- High-engagement outliers appear at mid-range durations.
- Very long videos do not consistently outperform shorter ones.
- Shorts within the optimal duration range maximize completion rates, which are a key signal for Shorts distribution algorithms.

🌟 Implication:

- Duration should be treated as a supporting, non-linear feature.
- Its effect likely depends on content type and context.

Key Insight: Shorts between 15–30 seconds show the highest median engagement.

Extremely short videos lack narrative depth, while longer videos experience viewer drop-off, reducing engagement signals.

EDA-Informed Feature Engineering Decisions

EDA revealed non-linear relationships and skewed engagement metrics, motivating the creation of rate-based features, log transformations, and time-based indicators such as peak upload hours.

EDA Summary

EDA reveals that engagement is driven by a combination of content length, category choice, publishing timing, and title characteristics. These patterns directly informed feature engineering and model design choices.

ML Approach & Model Training

Data Split & Pipelines

1 Target Encoding (Safe + Logged)

```
In [ ]: def encode_target(y: pd.Series):
        le = LabelEncoder()
        y_enc = le.fit_transform(y)

        mapping = dict(zip(le.classes_, le.transform(le.classes_)))
        print("Target Encoding Mapping:", mapping)

        return y_enc, le
```

2 Feature Split

```
In [ ]: def split_features_target(df: pd.DataFrame):
        drop_cols = [
            "video_id",
            "title",
            "engagement_rate",
            "performance_engagement_tertile"
        ]

        X = df.drop(columns=drop_cols, errors="ignore")
        y = df["performance_engagement_tertile"]

        return X, y
```

```
In [ ]: LEAKAGE_COLS = [
        "likes_per_sec", "comments_per_sec", "shares_per_sec",
        "log_likes", "log_comments", "log_shares", "log_views",
        "engagement_rate"
    ]
```

```
In [ ]: def split_features_target_new(df: pd.DataFrame):
        target_col = "performance_engagement_tertile"

        X = df.drop(
            columns=[
                "video_id",
                "title",
                target_col,
                *[c for c in LEAKAGE_COLS if c in df.columns]
            ],
```

```

        errors="ignore"
    )

    y = df[target_col]

    return X, y

```

3 Feature Type Detection

```

In [ ]: def detect_feature_types(X: pd.DataFrame):
        num_features = X.select_dtypes(include="number").columns.tolist()
        cat_features = X.select_dtypes(include="object").columns.tolist()

        # Treat upload_hour as categorical (business decision)
        if "upload_hour" in num_features:
            num_features.remove("upload_hour")
            cat_features.append("upload_hour")

        return num_features, cat_features

```

4 Preprocessor Builder

```

In [ ]: # value was 6 here its less than 20 so we can use OneHotEncoder for cat column
        df["category"].nunique()

```

Out[]: 6

```

In [ ]: def build_preprocessor(num_features, cat_features):
        return ColumnTransformer(
            transformers=[
                ("num", StandardScaler(), num_features),
                ("cat", OneHotEncoder(handle_unknown="ignore"), cat_features),
            ],
            remainder="drop"
        )

```

Model Training

5 Model Registry

```

In [ ]: def get_model_registry(random_state=42):
        return {
            # Linear
            "Logistic Regression": LogisticRegression(
                max_iter=1000,
                class_weight="balanced",
                random_state=random_state
            ),
            "ElasticNet Logistic": LogisticRegression(
                penalty="elasticnet",
                solver="saga",
                l1_ratio=0.5,
            )
        }

```

```

        max_iter=3000,
        class_weight="balanced",
        random_state=random_state
    ),

    # Tree-based

    "Random Forest": RandomForestClassifier(
        n_estimators=300,
        class_weight="balanced",
        random_state=random_state
    ),

    # Distance / Kernel

    "KNN": KNeighborsClassifier(),
    "SVM": SVC(
        probability=True,
        class_weight="balanced",
        random_state=random_state
    ),

    "Naive Bayes": GaussianNB(),

    # Boosting

    "XGBoost": XGBClassifier(
        objective="multi:softprob",
        eval_metric="mlogloss",
        n_estimators=300,
        max_depth=6,
        learning_rate=0.05,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=random_state
    ),
    "LightGBM": LGBMClassifier(
        n_estimators=300,
        learning_rate=0.05,
        num_leaves=31,
        subsample=0.8,
        colsample_bytree=0.8,
        class_weight="balanced",
        random_state=random_state
    ),
    "CatBoost": CatBoostClassifier(
        iterations=400,
        learning_rate=0.05,
        depth=6,
        loss_function="MultiClass",
        verbose=0,
        random_state=random_state
    ),
    "HistGradientBoosting": HistGradientBoostingClassifier(
        random_state=random_state

```

```
)  
  
}
```

6 Cross-Validation Engine (No Leakage)

```
In [ ]: def run_cross_validation(  
        X: pd.DataFrame,  
        y: np.ndarray,  
        preprocessor,  
        models: dict,  
        n_splits: int = 5,  
        random_state: int = 42  
    ):  
  
        """  
        Stratified K-Fold Cross-Validation for baseline model comparison.  
        Reports mean  $\pm$  std for Accuracy, Precision (macro), Recall (macro), F1 (macro).  
        """  
  
        Cv = StratifiedKFold(  
            n_splits=n_splits,  
            shuffle=True,  
            random_state=random_state  
        )  
  
        scoring = {  
            "accuracy": "accuracy",  
            "precision_macro": "precision_macro",  
            "recall_macro": "recall_macro",  
            "f1_macro": "f1_macro",  
        }  
  
        results = []  
  
        for name, model in models.items():  
            pipeline = Pipeline(  
                steps=[  
                    ("preprocessing", preprocessor),  
                    ("model", model),  
                ]  
            )  
  
            scores = cross_validate(  
                pipeline,  
                X,  
                y,  
                cv=Cv,  
                scoring=scoring,  
                n_jobs=-1,  
                return_train_score=False  
            )  
  
            results.append({  
                "Model": name,
```

```

        "Accuracy Mean": scores["test_accuracy"].mean(),
        "Accuracy Std": scores["test_accuracy"].std(),
        "Precision Macro Mean": scores["test_precision_macro"].mean(),
        "Precision Macro Std": scores["test_precision_macro"].std(),
        "Recall Macro Mean": scores["test_recall_macro"].mean(),
        "Recall Macro Std": scores["test_recall_macro"].std(),
        "F1 Macro Mean": scores["test_f1_macro"].mean(),
        "F1 Macro Std": scores["test_f1_macro"].std(),
    })

    return (
        pd.DataFrame(results)
        .sort_values("F1 Macro Mean", ascending=False)
        .reset_index(drop=True)
    )

```

Model Pipeline E2E Flow

```

In [ ]: # Split data
X, y = split_features_target_new(df)

# Encode target
y_encoded, label_encoder = encode_target(y)

# Detect features
num_features, cat_features = detect_feature_types(X)

# Build preprocessor
preprocessor = build_preprocessor(num_features, cat_features)

# Model registry
models = get_model_registry()

# Cross-validation
cv_results = run_cross_validation(
    X,
    y_encoded,
    preprocessor,
    models
)

cv_results

```

Target Encoding Mapping: {'High': np.int64(0), 'Low': np.int64(1), 'Medium': np.int64(2)}

Out[]:

	Model	Accuracy Mean	Accuracy Std	Precision Macro Mean	Precision Macro Std	Recall Macro Mean	Recall Macro Std	Macro F1
0	ElasticNet Logistic	0.943333	0.020000	0.945402	0.018887	0.942949	0.020203	0.943333
1	HistGradientBoosting	0.940000	0.029059	0.942454	0.028267	0.939808	0.029216	0.940000
2	XGBoost	0.933333	0.044721	0.934993	0.043791	0.933300	0.044759	0.933333
3	LightGBM	0.933333	0.045947	0.933794	0.046410	0.932790	0.046741	0.932790
4	CatBoost	0.923333	0.037417	0.926369	0.036618	0.923124	0.037250	0.923124
5	Logistic Regression	0.886667	0.035590	0.890161	0.034607	0.885581	0.035992	0.886667
6	SVM	0.840000	0.038873	0.844709	0.039787	0.839549	0.038997	0.840000
7	Random Forest	0.836667	0.038586	0.845727	0.039775	0.835547	0.039411	0.836667
8	KNN	0.660000	0.084063	0.659872	0.090303	0.659048	0.086404	0.656000
9	Naive Bayes	0.400000	0.039441	0.506449	0.069447	0.402047	0.038868	0.348889

Key insights :

F1-macro was selected as the primary model selection metric because all engagement classes (Low, Medium, High) are equally important from a business standpoint. Accuracy alone can mask poor minority-class performance, whereas F1-macro enforces balanced classification quality.

📊 Cross-Validation Model Comparison (5-Fold CV)			
Model	Accuracy (Mean ± Std)	F1-Macro (Mean ± Std)	Key Takeaway
ElasticNet Logistic	0.943 ± 0.020	0.943 ± 0.020	Best overall, stable & interpretable
HistGradientBoosting	0.940 ± 0.029	0.940 ± 0.029	Strong non-linear learner
XGBoost	0.933 ± 0.045	0.933 ± 0.045	High variance on small dataset
LightGBM	0.933 ± 0.047	0.933 ± 0.047	Similar performance, less stable
CatBoost	0.923 ± 0.037	0.924 ± 0.037	Good but slightly weaker
Random Forest	0.837 ± 0.040	0.837 ± 0.040	Underperformed
SVM	0.841 ± 0.039	0.841 ± 0.039	Lower recall balance
KNN	0.660 ± 0.092	0.656 ± 0.092	Highly unstable
Naive Bayes	0.400 ± 0.039	0.349 ± 0.049	Poor fit

Baseline Model Experiments & Comparison (Cross-Validation Results)

```
In [ ]: def set_plot_style():
    sns.set_theme(style="whitegrid")
    plt.rcParams.update({
        "figure.dpi": 120,
        "axes.titlesize": 12,
        "axes.labelsize": 10,
        "xtick.labelsize": 9,
        "ytick.labelsize": 9,
    })
```

Bar Plot – F1 Macro Mean (Baseline Ranking)

Purpose: Quickly compare average performance across models.

```
In [ ]: def plot_f1_macro_bar(
    cv_results_df: pd.DataFrame,
    save_path: Optional[str] = None
):
    required_cols = {"Model", "F1 Macro Mean"}
    if not required_cols.issubset(cv_results_df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    plt.figure(figsize=(10, 10))
    sns.barplot(
        data=cv_results_df.sort_values("F1 Macro Mean", ascending=False),
        x="Model",
        y="F1 Macro Mean",
        palette="viridis"
    )

    plt.title("Model Comparison using F1 Macro (Cross-Validation)")
    plt.xlabel("Model")
    plt.ylabel("F1 Macro Mean")
    plt.xticks(rotation=30)

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()
```

Error Bar Plot – F1 Macro Mean $\pm$ Std (Stability)

Purpose: Understand model stability / variance across folds.

```
In [ ]: def plot_f1_macro_with_std(
    cv_results_df: pd.DataFrame,
    save_path: Optional[str] = None
):
    required_cols = {
        "Model",
        "F1 Macro Mean",
```

```

        "F1 Macro Std"
    }
    if not required_cols.issubset(cv_results_df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    df_sorted = cv_results_df.sort_values(
        by="F1 Macro Mean", ascending=False
    )

    plt.figure(figsize=(10, 10))
    plt.bar(
        df_sorted["Model"],
        df_sorted["F1 Macro Mean"],
        yerr=df_sorted["F1 Macro Std"],
        capsize=5
    )

    plt.title("F1 Macro Mean  $\pm$  Std (Cross-Validation)")
    plt.xlabel("Model")
    plt.ylabel("F1 Macro")
    plt.xticks(rotation=30)

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

```

3 Heatmap – Multi-Metric Comparison

Purpose: Reveal trade-offs between Accuracy, Precision, Recall, F1

```

In [ ]: def plot_model_performance_heatmap(
    cv_results_df: pd.DataFrame,
    metrics_cols: Optional[List[str]] = None,
    save_path: Optional[str] = None
):
    if metrics_cols is None:
        metrics_cols = [
            "Accuracy Mean",
            "Precision Macro Mean",
            "Recall Macro Mean",
            "F1 Macro Mean"
        ]

    required_cols = {"Model", *metrics_cols}
    if not required_cols.issubset(cv_results_df.columns):
        raise ValueError(f"Missing columns: {required_cols}")

    plt.figure(figsize=(8, 10))

    sns.heatmap(
        cv_results_df
        .set_index("Model")[metrics_cols],
        annot=True,
        fmt=".3f",

```

```

        cmap="YlGnBu",
        linewidths=0.5
    )

    plt.title("Cross-Validated Model Performance Heatmap")

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

```

```


In [ ]: def baseline_Model_Comparison_Plots(
        cv_results_df: pd.DataFrame
    ):
        set_plot_style()
        plot_f1_macro_bar(cv_results_df)
        print()
        print(" *" * 70)
        plot_f1_macro_with_std(cv_results_df)
        print()
        print(" *" * 70)
        print()
        plot_model_performance_heatmap(cv_results_df)
        print(" *" * 70)


```


```

In [ ]: baseline_Model_Comparison_Plots(cv_results)

```







Key insights :

Final Selected Model: ElasticNet Logistic Regression Primary Selection Metric: F1-macro

Hyperparameter Tuning and Model Optimization

Top 5 Model Selection as per baseline performane

```
In [ ]: #models_to_tune = list(models.keys())

models_to_tune = cv_results.head(5)["Model"].tolist()
```

```
In [ ]: models_to_tune
```

```
Out[ ]: ['ElasticNet Logistic',
        'HistGradientBoosting',
        'XGBoost',
        'LightGBM',
        'CatBoost']
```

```
In [ ]: def filter_model_registry(model_registry: dict, allowed_models: list) -> dict:
        return {
            name: model
            for name, model in model_registry.items()
            if name in allowed_models
        }
```

```
In [ ]: selectedModels = filter_model_registry(models, models_to_tune)
print(selectedModels)
```

```
{'ElasticNet Logistic': LogisticRegression(class_weight='balanced', l1_ratio=0.5, max_iter=3000,
                                           penalty='elasticnet', random_state=42, solver='saga'), 'XGBoost':
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.8, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric='mlogloss',
              feature_types=None, feature_weights=None, gamma=None,
              grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.05, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=6, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=300, n_jobs=None,
              num_parallel_tree=None, ...), 'LightGBM': LGBMClassifier(class_weight
='balanced', colsample_bytree=0.8,
              learning_rate=0.05, n_estimators=300, random_state=42,
              subsample=0.8), 'CatBoost': <catboost.core.CatBoostClassifier object
at 0x7bc09dcd4050>, 'HistGradientBoosting': HistGradientBoostingClassifier(random_state=42)}
```

1 Define parameter grids

```
In [ ]: PARAM_GRIDS = {

    # ----- Linear -----
    "Logistic Regression": {
        "model__C": [0.01, 0.1, 1, 10],
        "model__penalty": ["l2"]
    },

    "ElasticNet Logistic": {
        "model__C": [0.01, 0.1, 1, 10],
```

```

        "model__l1_ratio": [0.2, 0.5, 0.8]
    },

    # ----- Tree -----
    "Random Forest": {
        "model__n_estimators": [200, 400, 600],
        "model__max_depth": [None, 10, 20],
        "model__min_samples_split": [2, 5],
        "model__min_samples_leaf": [1, 2]
    },

    # ----- Distance / Kernel -----
    "KNN": {
        "model__n_neighbors": [5, 10, 20, 30],
        "model__weights": ["uniform", "distance"],
        "model__p": [1, 2] # Manhattan vs Euclidean
    },

    "SVM": {
        "model__C": [0.1, 1, 10],
        "model__gamma": ["scale", "auto"],
        "model__kernel": ["rbf"]
    },

    # ----- Probabilistic -----
    "Naive Bayes": {
        "model__var_smoothing": [1e-9, 1e-8, 1e-7]
    },

    # ----- Boosting -----
    "XGBoost": {
        "model__n_estimators": [300, 600],
        "model__learning_rate": [0.03, 0.05, 0.1],
        "model__max_depth": [3, 6],
        "model__subsample": [0.8, 1.0],
        "model__colsample_bytree": [0.8, 1.0]
    },

    "LightGBM": {
        "model__n_estimators": [300, 600],
        "model__learning_rate": [0.03, 0.05, 0.1],
        "model__num_leaves": [31, 63],
        "model__subsample": [0.8, 1.0],
        "model__colsample_bytree": [0.8, 1.0]
    },

    "CatBoost": {
        "model__iterations": [300, 600],
        "model__learning_rate": [0.03, 0.05, 0.1],
        "model__depth": [4, 6, 8]
    },

    "HistGradientBoosting": {
        "model__learning_rate": [0.03, 0.05, 0.1],
        "model__max_depth": [6, 8, 10],
        "model__max_iter": [200, 400]
    }

```

```
}
}
```

2 RandomizedSearch and cvsearch for ALL models

```
In [ ]: CV = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
In [ ]: def run_randomized_tuning(
    X, y, preprocessor, models, param_grids, n_iter=20
):
    results = []

    for name, model in models.items():
        print(f"\n🔍 RandomizedSearch → {name}")

        pipeline = Pipeline([
            ("preprocessing", preprocessor),
            ("model", model)
        ])

        search = RandomizedSearchCV(
            pipeline,
            param_distributions=param_grids[name],
            n_iter=n_iter,
            scoring="f1_macro",
            cv=CV,
            n_jobs=-1,
            random_state=42
        )

        search.fit(X, y)

        results.append({
            "Model": name,
            "Search": "Randomized",
            "Best F1": search.best_score_,
            "Best Params": search.best_params_
        })

    return pd.DataFrame(results).sort_values(
        "Best F1", ascending=False
    )
```

```
In [ ]: random_tuning_results = run_randomized_tuning(
    X,
    y_encoded,
    preprocessor,
    selectedModels,
    PARAM_GRIDS
)

random_tuning_results
```

RandomizedSearch → ElasticNet Logistic

RandomizedSearch → XGBoost

RandomizedSearch → LightGBM

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000051 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 484

[LightGBM] [Info] Number of data points in the train set: 300, number of used features: 14

[LightGBM] [Info] Start training from score -1.098612

[LightGBM] [Info] Start training from score -1.098612

[LightGBM] [Info] Start training from score -1.098612

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

 RandomizedSearch → CatBoost

 RandomizedSearch → HistGradientBoosting

Out[]:

	Model	Search	Best F1	Best Params
0	ElasticNet Logistic	Randomized	0.976658	{'model__l1_ratio': 0.8, 'model__C': 1}
4	HistGradientBoosting	Randomized	0.946702	{'model__max_iter': 400, 'model__max_depth': 8...
1	XGBoost	Randomized	0.943302	{'model__subsample': 0.8, 'model__n_estimators...
2	LightGBM	Randomized	0.943038	{'model__subsample': 0.8, 'model__num_leaves':...
3	CatBoost	Randomized	0.930326	{'model__learning_rate': 0.1, 'model__iteratio...

```
In [ ]: def run_gridSearch_cv_tuning(
        X, y, preprocessor, models, param_grids
    ):
        results = []

        for name, model in models.items():
            print(f"\n🔍 GridSearch → {name}")

            pipeline = Pipeline([
                ("preprocessing", preprocessor),
                ("model", model)
            ])

            search = GridSearchCV(
                pipeline,
                param_grid=param_grids[name],
                scoring="f1_macro",
```

```

        cv=CV,
        n_jobs=-1
    )

    search.fit(X, y)

    results.append({
        "Model": name,
        "Search": "Grid",
        "Best F1": search.best_score_,
        "Best Params": search.best_params_
    })

    return pd.DataFrame(results).sort_values(
        "Best F1", ascending=False
    )

```

```

In [ ]: gridSearchcv_tuning_results = run_gridSearch_cv_tuning(
        X,
        y_encoded,
        preprocessor,
        selectedModels,
        PARAM_GRIDS
    )

gridSearchcv_tuning_results

```

 GridSearch → ElasticNet Logistic

 GridSearch → XGBoost

 GridSearch → LightGBM

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000076 seconds.

You can set `force\_col\_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 484

[LightGBM] [Info] Number of data points in the train set: 300, number of used features: 14

[LightGBM] [Info] Start training from score -1.098612

[LightGBM] [Info] Start training from score -1.098612

[LightGBM] [Info] Start training from score -1.098612

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

 GridSearch → CatBoost

 GridSearch → HistGradientBoosting

Out []:

	Model	Search	Best F1	Best Params
0	ElasticNet Logistic	Grid	0.976658	{'model__C': 1, 'model__l1_ratio': 0.8}
4	HistGradientBoosting	Grid	0.946702	{'model__learning_rate': 0.05, 'model__max_dep...
1	XGBoost	Grid	0.943375	{'model__colsample_bytree': 1.0, 'model__learn...
2	LightGBM	Grid	0.943038	{'model__colsample_bytree': 1.0, 'model__learn...
3	CatBoost	Grid	0.930326	{'model__depth': 8, 'model__iterations': 300, ...

Tuned Model Performance

In []: random_tuning_results

Out[]:	Model	Search	Best F1	Best Params
0	ElasticNet Logistic	Randomized	0.976658	{'model__l1_ratio': 0.8, 'model__C': 1}
4	HistGradientBoosting	Randomized	0.946702	{'model__max_iter': 400, 'model__max_depth': 8...
1	XGBoost	Randomized	0.943302	{'model__subsample': 0.8, 'model__n_estimators...
2	LightGBM	Randomized	0.943038	{'model__subsample': 0.8, 'model__num_leaves':...
3	CatBoost	Randomized	0.930326	{'model__learning_rate': 0.1, 'model__iteratio...

In []: `gridSearchcv_tuning_results`

Out[]:	Model	Search	Best F1	Best Params
0	ElasticNet Logistic	Grid	0.976658	{'model__C': 1, 'model__l1_ratio': 0.8}
4	HistGradientBoosting	Grid	0.946702	{'model__learning_rate': 0.05, 'model__max_dep...
1	XGBoost	Grid	0.943375	{'model__colsample_bytree': 1.0, 'model__learn...
2	LightGBM	Grid	0.943038	{'model__colsample_bytree': 1.0, 'model__learn...
3	CatBoost	Grid	0.930326	{'model__depth': 8, 'model__iterations': 300, ...

Final Model Production Ready Comparision

```
In [ ]: def set_plot_style():
sns.set_theme(style="whitegrid")
plt.rcParams.update({
    "figure.dpi": 120,
    "axes.titlesize": 12,
    "axes.labelsize": 10,
    "xtick.labelsize": 9,
    "ytick.labelsize": 9,
})
```

```
In [ ]: def build_tuning_comparison_df(
    baseline_df: pd.DataFrame,
    grid_df: pd.DataFrame,
    random_df: pd.DataFrame,
    sort_by: str = "GridSearch F1"
) -> pd.DataFrame:
    baseline = baseline_df[["Model", "F1 Macro Mean"]].rename(
        columns={"F1 Macro Mean": "Baseline F1"}
    )

    grid = grid_df[["Model", "Best F1"]].rename(
        columns={"Best F1": "GridSearch F1"}
```

```

)

random = random_df[["Model", "Best F1"]].rename(
    columns={"Best F1": "RandomSearch F1"}
)

comparison_df = (
    baseline
    .merge(grid, on="Model", how="inner")
    .merge(random, on="Model", how="inner")
    .sort_values(by=sort_by, ascending=False)
    .reset_index(drop=True)
)

return comparison_df

```

Baseline vs Grid vs Random (Bar Chart)

```

In [ ]: def plot_f1_comparison_bar(
    comparison_df: pd.DataFrame,
    save_path: Optional[str] = None
):
    comparison_df.set_index("Model").plot(
        kind="bar",
        figsize=(12, 6)
    )

    plt.title("Baseline vs GridSearch vs RandomizedSearch (F1-Macro)")
    plt.ylabel("F1-Macro")
    plt.xlabel("Model")
    plt.xticks(rotation=30)
    plt.legend(title="Evaluation Type")
    plt.grid(axis="y", linestyle="--", alpha=0.6)

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

```

F1 improvement

```

In [ ]: def plot_f1_improvement(
    comparison_df: pd.DataFrame,
    save_path: Optional[str] = None
):
    improvement_df = comparison_df.copy()

    improvement_df["Grid Improvement"] = (
        improvement_df["GridSearch F1"] - improvement_df["Baseline F1"]
    )
    improvement_df["Random Improvement"] = (
        improvement_df["RandomSearch F1"] - improvement_df["Baseline F1"]
    )

    improvement_df.set_index("Model")[

```

```

        ["Grid Improvement", "Random Improvement"]
    ].plot(
        kind="bar",
        figsize=(12, 5)
    )

    plt.title("F1-Macro Improvement from Hyperparameter Tuning")
    plt.ylabel("Δ F1-Macro")
    plt.xlabel("Model")
    plt.xticks(rotation=30)
    plt.axhline(0, color="black", linewidth=1)
    plt.grid(axis="y", linestyle="--", alpha=0.6)

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

```

Heatmap (Compact Decision View)

```

In [ ]: def plot_f1_heatmap(
    comparison_df: pd.DataFrame,
    save_path: Optional[str] = None
):
    plt.figure(figsize=(10, 5))
    sns.heatmap(
        comparison_df.set_index("Model"),
        annot=True,
        cmap="YlGnBu",
        fmt=".3f"
    )

    plt.title("Model Performance Comparison (F1-Macro)")

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()

```

Trend Across Tuning Stages (Line Plot)

```

In [ ]: def plot_f1_trend(
    comparison_df: pd.DataFrame,
    save_path: Optional[str] = None
):
    stages = ["Baseline F1", "GridSearch F1", "RandomSearch F1"]

    plt.figure(figsize=(12, 6))

    for _, row in comparison_df.iterrows():
        plt.plot(
            stages,
            [row[s] for s in stages],
            marker="o",
            label=row["Model"]
        )

```

```

plt.title("F1-Macro Trend Across Model Tuning Stages")
plt.ylabel("F1-Macro Score")
plt.xlabel("Model Training Stage")
plt.grid(True, linestyle="--", alpha=0.6)
plt.legend(title="Model", bbox_to_anchor=(1.05, 1), loc="upper left")
plt.tight_layout()

if save_path:
    plt.savefig(save_path, bbox_inches="tight")
plt.show()

```

```

In [ ]: def generate_full_tuning_comparison_report(
        baseline_df: pd.DataFrame,
        grid_df: pd.DataFrame,
        random_df: pd.DataFrame,
        output_dir: Optional[str] = None
    ):
        set_plot_style()

        comparison_df = build_tuning_comparison_df(
            baseline_df,
            grid_df,
            random_df
        )

        plot_f1_comparison_bar(
            comparison_df,
            save_path=f"{output_dir}/f1_comparison_bar.png" if output_dir else None
        )

        plot_f1_improvement(
            comparison_df,
            save_path=f"{output_dir}/f1_improvement.png" if output_dir else None
        )

        plot_f1_heatmap(
            comparison_df,
            save_path=f"{output_dir}/f1_heatmap.png" if output_dir else None
        )

        plot_f1_trend(
            comparison_df,
            save_path=f"{output_dir}/f1_trend.png" if output_dir else None
        )

        print("✅ Hyperparameter tuning comparison report generated")

        return comparison_df

```



Final Model Selection



```

In [ ]: final_comparison_df = generate_full_tuning_comparison_report(
        baseline_df=cv_results,
        grid_df=gridSearchcv_tuning_results,
        random_df=random_tuning_results,

```

```
output_dir="" # optional
)
```





✓ Hyperparameter tuning comparison report generated

Final Model Selection Criteria

Define Selection Factors & Rationale

```
In [ ]: SELECTION_FACTORS = {  
    "Training speed": "Faster retraining & iteration",  
    "Parameter sensitivity": "Robustness to tuning",  
    "Stability": "Low variance across CV folds",  
    "Production adoption": "Ease of deployment & explainability",  
}
```

Define Top-3 Models With Strengths & Trade-offs

```
In [ ]: TOP3_MODEL_PROFILE = {  
    "ElasticNet Logistic": {  
        "strengths": [  

```

```

        "Fast training and inference",
        "Highly stable across folds",
        "Interpretable coefficients",
        "Minimal hyperparameter sensitivity",
        "Best tuned F1-macro"
    ],
    "trade_offs": [
        "Assumes linear decision boundaries"
    ]
},
"XGBoost": {
    "strengths": [
        "Strong non-linear modeling",
        "Handles feature interactions well"
    ],
    "trade_offs": [
        "Slower training",
        "Highly sensitive to hyperparameters",
        "Higher operational complexity"
    ]
},
"HistGradientBoosting": {
    "strengths": [
        "Balanced performance",
        "Efficient histogram-based training",
        "Low tuning complexity"
    ],
    "trade_offs": [
        "Smaller ecosystem",
        "Limited production tooling"
    ]
}
}

```

Score Models on Production Factors

```

In [ ]: #Scale: 1 (poor) → 5 (excellent)

TOP3_FACTOR_SCORES = {
    "ElasticNet Logistic": {
        "Training speed": 5,
        "Parameter sensitivity": 5,
        "Stability": 5,
        "Production adoption": 5,
    },
    "XGBoost": {
        "Training speed": 3,
        "Parameter sensitivity": 2,
        "Stability": 3,
        "Production adoption": 4,
    },
    "HistGradientBoosting": {
        "Training speed": 4,
        "Parameter sensitivity": 4,
        "Stability": 4,
    }
}

```

```

        "Production adoption": 3,
    },
}

```

```

In [ ]: def select_final_model_beyond_accuracy(
    factor_scores: dict,
    model_profiles: dict,
    weights: dict | None = None,
    verbose: bool = True
):
    """
    Select final model using production-oriented factors
    instead of accuracy alone.
    """

    factors = list(next(iter(factor_scores.values())).keys())

    if weights is None:
        weights = {factor: 1.0 for factor in factors}

    final_scores = {}

    for model, scores in factor_scores.items():
        final_scores[model] = sum(
            scores[factor] * weights[factor]
            for factor in factors
        )

    selected_model = max(final_scores, key=final_scores.get)

    if verbose:
        print("\n🚀 Model Selection Criteria (Beyond Accuracy)")
        for factor, reason in SELECTION_FACTORS.items():
            print(f"- {factor}: {reason}")

        print("\n🔍 Top-3 Model Comparison\n")
        for model, score in sorted(
            final_scores.items(),
            key=lambda x: x[1],
            reverse=True
        ):
            print(f"{model:25s} → Total Score: {score:.2f}")

        print("\n🏆 Final Selected Model:")
        print(f"➡ {selected_model}\n")

        print("✅ Why this model was chosen:")
        for reason in model_profiles[selected_model]["strengths"]:
            print(f"    • {reason}")

        print("\n⚠️ Trade-offs considered:")
        for trade in model_profiles[selected_model]["trade_offs"]:
            print(f"    • {trade}")

        print(
            "\n📌 Conclusion:\n"

```

```

        "The marginal performance gains of more complex boosting models "
        "did not justify their additional tuning, training time, and "
        "operational overhead in a production setting."
    )

    return selected_model, final_scores

```

```

In [ ]: FACTOR_WEIGHTS = {
        "Training speed": 1.0,
        "Parameter sensitivity": 1.5,
        "Stability": 2.0,
        "Production adoption": 2.0,
    }

```

```

In [ ]: final_model_name, score_table = select_final_model_beyond_accuracy(
        factor_scores=TOP3_FACTOR_SCORES,
        model_profiles=TOP3_MODEL_PROFILE,
        weights=FACTOR_WEIGHTS
    )

```

🚩 Model Selection Criteria (Beyond Accuracy)

- Training speed: Faster retraining & iteration
- Parameter sensitivity: Robustness to tuning
- Stability: Low variance across CV folds
- Production adoption: Ease of deployment & explainability

🔍 Top-3 Model Comparison

ElasticNet Logistic	→ Total Score: 32.50
HistGradientBoosting	→ Total Score: 24.00
XGBoost	→ Total Score: 20.00

🏆 Final Selected Model:

➡ ElasticNet Logistic

✅ Why this model was chosen:

- Fast training and inference
- Highly stable across folds
- Interpretable coefficients
- Minimal hyperparameter sensitivity
- Best tuned F1-macro

⚠️ Trade-offs considered:

- Assumes linear decision boundaries

📌 Conclusion:

The marginal performance gains of more complex boosting models did not justify their additional tuning, training time, and operational overhead in a production setting.

FINAL MODEL DECLARATION :

Final Selected Model: ElasticNet Logistic Regression Primary Selection Metric: F1-macro

STABILITY EXPLANATION:

ElasticNet Logistic Regression demonstrates not only the highest mean performance but also the lowest standard deviation, indicating stable generalization across folds.

REJECTION LOGIC FOR STRONG MODELS

Although HistGradientBoosting, XGBoost, and LightGBM achieved competitive performance, their higher variance and reduced interpretability introduce greater overfitting risk given the limited dataset size (n=300).

Why ElasticNet Logistic Regression?

ElasticNet regularization combines L1 (feature selection) and L2 (coefficient stability) penalties. This approach is particularly effective when features are correlated, as is common in engagement-related metrics, helping prevent overfitting while maintaining strong predictive performance.

Model Success Criteria

The model is considered successful if it achieves balanced F1-macro performance above 0.90 while maintaining interpretability for actionable content strategy decisions.

Note:

“The final model was selected using production-oriented criteria—training speed, tuning robustness, stability, and deployment suitability—rather than accuracy alone, leading to the selection of ElasticNet Logistic Regression.”

Final Evaluation

Train and evaluate final model

```
In [ ]: def train_and_evaluate_final_model(
        X: pd.DataFrame,
        y: np.ndarray,
        preprocessor,
        model,
        best_params: Optional[Dict[str, str]] = None,
        test_size: float = 0.2,
        random_state: int = 42,
        verbose: bool = True,
    ):
        """
        Train final model with tuned hyperparameters and evaluate on a held-out test set.

        Returns trained pipeline and structured evaluation outputs.
        """

        # -----
        # Train / Test split
```

```

# -----
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=test_size,
    stratify=y,
    random_state=random_state,
)

# -----
# Build pipeline
# -----
pipeline = Pipeline(
    steps=[
        ("preprocessing", preprocessor),
        ("model", model),
    ]
)

# -----
# Apply tuned hyperparameters
# -----
if best_params:
    pipeline.set_params(**best_params)

# -----
# Fit model
# -----
pipeline.fit(X_train, y_train)

# -----
# Predictions
# -----
y_pred = pipeline.predict(X_test)

# -----
# Probabilities + ROC-AUC
# -----
roc_auc = np.nan
y_proba = None

if hasattr(pipeline.named_steps["model"], "predict_proba"):
    y_proba = pipeline.predict_proba(X_test)
    roc_auc = roc_auc_score(
        y_test,
        y_proba,
        multi_class="ovr",
        average="macro",
    )

# -----
# Metrics
# -----
metrics = {
    "Accuracy": accuracy_score(y_test, y_pred),
    "F1 Macro": f1_score(y_test, y_pred, average="macro"),
}

```

```

        "ROC-AUC OvR": roc_auc,
    }

    cm = confusion_matrix(y_test, y_pred)
    report_dict = classification_report(
        y_test, y_pred, output_dict=True
    )
    report_df = pd.DataFrame(report_dict).T

    # -----
    # Optional console output
    # -----
    if verbose:
        print("\n❖ Final Model Evaluation (Held-out Test Set)\n")
        for k, v in metrics.items():
            print(f"{k:<14}: {v:.4f}")
        print("\nClassification Report:\n")
        print(classification_report(y_test, y_pred))

    # -----
    # Return structured results
    # -----
    return {
        "pipeline": pipeline,
        "X_train": X_train,
        "X_test": X_test,
        "y_train": y_train,
        "y_test": y_test,
        "y_pred": y_pred,
        "y_proba": y_proba,
        "confusion_matrix": cm,
        "classification_report": report_df,
        "metrics": metrics,
    }

```

Select Best Model params

```

In [ ]: def get_best_params_by_model(
        tuning_df: pd.DataFrame,
        model_name: str
    ) -> dict:
        row = tuning_df.loc[
            tuning_df["Model"] == model_name,
            "Best Params"
        ]

        if row.empty:
            raise ValueError(f"Model '{model_name}' not found")

        return row.iloc[0]

```

Evalute Final Report

```
In [ ]: final_model = models[final_model_name]
print("Final Model:" , final_model_name)
best_params = get_best_params_by_model(
    gridSearchcv_tuning_results,
    final_model_name
)
print()
print("Best Params:")
print( best_params)

result = train_and_evaluate_final_model(
    X=X,
    y=y_encoded,
    preprocessor=preprocessor,
    model=final_model,
    best_params=best_params
)

result["metrics"]
result["confusion_matrix"]
result["classification_report"]
```

Final Model: ElasticNet Logistic

Best Params:

{'model__C': 1, 'model__l1_ratio': 0.8}

❖ Final Model Evaluation (Held-out Test Set)

Accuracy : 0.9833
F1 Macro : 0.9833
ROC-AUC OvR : 0.9992

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.95	0.97	20
1	1.00	1.00	1.00	20
2	0.95	1.00	0.98	20
accuracy			0.98	60
macro avg	0.98	0.98	0.98	60
weighted avg	0.98	0.98	0.98	60

		precision	recall	f1-score	support
	0	1.000000	0.950000	0.974359	20.000000
	1	1.000000	1.000000	1.000000	20.000000
	2	0.952381	1.000000	0.975610	20.000000
	accuracy	0.983333	0.983333	0.983333	0.983333
	macro avg	0.984127	0.983333	0.983323	60.000000
	weighted avg	0.984127	0.983333	0.983323	60.000000

Confusion Matrix Interpretation

```
In [ ]: def plot_confusion_matrix(
    y_true,
    y_pred,
    class_labels,
    title: str = "Confusion Matrix",
    cmap: str = "Blues",
    save_path: Optional[str] = None,
):
    cm = confusion_matrix(y_true, y_pred)

    plt.figure(figsize=(6, 5))
    sns.heatmap(
        cm,
        annot=True,
        fmt="d",
        cmap=cmap,
        xticklabels=class_labels,
        yticklabels=class_labels,
    )

    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.title(title)

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")
    plt.show()
```

Final Test Metrics Table + Bar Chart

```
In [ ]: def build_test_metrics_table(
    y_true,
    y_pred,
    y_proba=None,
):
    metrics = {
        "Accuracy": accuracy_score(y_true, y_pred),
        "F1 Macro": f1_score(y_true, y_pred, average="macro"),
```

```

}

if y_proba is not None:
    metrics["ROC-AUC (OvR)"] = roc_auc_score(
        y_true,
        y_proba,
        multi_class="ovr",
        average="macro",
    )

return pd.DataFrame(
    {"Metric": metrics.keys(), "Value": metrics.values()}
)

```

Metrics Bar Chart

```

In [ ]: def plot_metrics_bar(
    metrics_df: pd.DataFrame,
    title: str,
    save_path: str | None = None,
):
    plt.figure(figsize=(6, 4))

    bars = plt.bar(
        metrics_df["Metric"],
        metrics_df["Value"],
    )

    plt.title(title)
    plt.ylabel("Score")
    plt.ylim(0, 1.05)
    plt.grid(axis="y", linestyle="--", alpha=0.6)

    for bar in bars:
        height = bar.get_height()
        plt.text(
            bar.get_x() + bar.get_width() / 2,
            height + 0.01,
            f"{height:.3f}",
            ha="center",
            va="bottom",
            fontsize=10,
        )

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")

    plt.show()

```

Build Train/Test Metrics

```

In [ ]: def build_train_test_comparison(
    y_train,
    y_train_pred,

```

```

y_test,
y_test_pred,
):
    return pd.DataFrame({
        "Dataset": ["Train", "Test"],
        "Accuracy": [
            accuracy_score(y_train, y_train_pred),
            accuracy_score(y_test, y_test_pred),
        ],
        "F1 Macro": [
            f1_score(y_train, y_train_pred, average="macro"),
            f1_score(y_test, y_test_pred, average="macro"),
        ],
    })

```

```

In [ ]: def plot_train_test_comparison(
    comparison_df: pd.DataFrame,
    title: str = "Train vs Test Performance",
    save_path: str | None = None,
):
    melted = comparison_df.melt(
        id_vars="Dataset",
        var_name="Metric",
        value_name="Score"
    )

    plt.figure(figsize=(7, 5))
    sns.barplot(
        data=melted,
        x="Metric",
        y="Score",
        hue="Dataset",
    )

    plt.ylim(0, 1.05)
    plt.title(title)
    plt.grid(axis="y", linestyle="--", alpha=0.6)

    if save_path:
        plt.savefig(save_path, bbox_inches="tight")

    plt.show()

```

Final Report New Model validation

```

In [ ]: print("Confusion Matrix")
plot_confusion_matrix(
    result["y_test"],
    result["y_pred"],
    class_labels=label_encoder.classes_,
    title="Confusion Matrix - ElasticNet Logistic (Final Test Set)"
)

final_test_metrics = build_test_metrics_table(
    result["y_test"],

```


```

    result["y_pred"],
    result["y_proba"]
)
print("Final Model Performance")


plot_metrics_bar(
    final_test_metrics,
    title="Final Model Performance on Test Set"
)

```

Confusion Matrix



Final Model Performance



Confusion Matrix Interpretation

Key Observations: • Medium engagement Shorts are most frequently confused with High • Recall for High engagement remains strong, which is critical for identifying viral candidates

Interpretation of Medium Engagement Class


Medium engagement Shorts represent optimization opportunities. These videos often require small adjustments in duration, title structure, or upload timing to transition into the High engagement category.

Train vs Test comparison

```
In [ ]: # Predictions on train
y_train_pred = result["pipeline"].predict(result["X_train"]) # example

train_test_df = build_train_test_comparison(
    result["y_train"],      # true train labels (240)
    y_train_pred,          # train predictions (240)
    result["y_test"],       # true test labels (60)
    result["y_pred"],       # test predictions (60)
)
print("Train vs Test Performance")
plot_train_test_comparison(
    train_test_df,
    title="Train vs Test Performance - Final Model"
)
```

Train vs Test Performance



Train vs Test Performance Analysis

The test set performance closely aligns with cross-validation results, indicating minimal overfitting and stable generalization behavior. This suggests the selected model is robust and suitable for deployment.

ROC-AUC (One-vs-Rest)

```
In [ ]: def plot_multiclass_roc_curve(
    pipeline,
    X_test,
    y_test,
    class_labels: Optional[list] = None,
    title: str = "Multiclass ROC Curve (One-vs-Rest)",
    figsize: tuple = (7, 6),
    save_path: Optional[str] = None,
):
    """
    Plot One-vs-Rest ROC curves for a multiclass classifier.

    Parameters
    -----
    pipeline : sklearn Pipeline
        Trained pipeline with a classifier supporting predict_proba.
    X_test : pd.DataFrame or np.ndarray
```

```

    Test features.
y_test : array-like
    True test labels (encoded).
class_labels : list, optional
    Class labels in correct order. If None, inferred from y_test.
title : str
    Plot title.
figsize : tuple
    Figure size.
save_path : str, optional
    If provided, saves the figure to this path.

Returns
-----
dict
    Per-class AUC scores.
"""

# -----
# Validation
# -----
if not hasattr(pipeline.named_steps["model"], "predict_proba"):
    raise ValueError("Model does not support predict_proba; ROC cannot be computed")

y_test = np.asarray(y_test)

if class_labels is None:
    class_labels = np.unique(y_test)

n_classes = len(class_labels)

# -----
# Binarize Labels
# -----
y_test_bin = label_binarize(y_test, classes=class_labels)

# -----
# Predict probabilities
# -----
y_proba = pipeline.predict_proba(X_test)

if y_proba.shape[1] != n_classes:
    raise ValueError(
        f"Mismatch: {n_classes} classes but predict_proba returned {y_proba.shape[1]} classes"
    )

# -----
# Plot
# -----
plt.figure(figsize=figsize)
auc_scores = {}

for i, cls in enumerate(class_labels):
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_proba[:, i])
    roc_auc = auc(fpr, tpr)
    auc_scores[cls] = roc_auc

```

```

plt.plot(
    fpr,
    tpr,
    linewidth=2,
    label=f"Class {cls} (AUC = {roc_auc:.3f})",
)

# Random baseline
plt.plot([0, 1], [0, 1], linestyle="--", linewidth=1)

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title(title)
plt.legend(loc="lower right")
plt.grid(alpha=0.3)

if save_path:
    plt.savefig(save_path, bbox_inches="tight")

plt.show()

return auc_scores


```

```

In [ ]: auc_scores = plot_multiclass_roc_curve(
    pipeline=result["pipeline"],
    X_test=result["X_test"],
    y_test=result["y_test"],
    title="Multiclass ROC Curve - Final Model",
)

print("Per-class AUC:")
for cls, score in auc_scores.items():
    print(f"Class {cls}: {score:.4f}")

```



Per-class AUC:
Class 0: 1.0000
Class 1: 1.0000
Class 2: 0.9975

Key insights:

ROC-AUC (One-vs-Rest) was evaluated as a secondary metric to assess the model's ability to separate engagement classes. The selected model demonstrates strong class separability, enabling reliable ranking of Shorts by viral potential.

Assumptions & Limitations

- The dataset consists of only 300 YouTube Shorts, which may limit the generalizability of the results.
- Engagement metrics are observational and do not imply causality.
- Temporal trends and algorithmic changes on YouTube are not explicitly modeled.
- The findings should be validated on larger, real-world datasets before production deployment.

Model Explainability & Business Insights

Model Explainability : Coefficient Analysis - ElasticNet Logistic Regression

```
In [ ]: model = result["pipeline"].named_steps["model"]

coef_df = pd.DataFrame(
    {
        "feature": result["pipeline"]
            .named_steps["preprocessing"]
            .get_feature_names_out(),
        "coefficient": model.coef_[0],
    }
)

coef_df["abs_coefficient"] = coef_df["coefficient"].abs()
coef_df = coef_df.sort_values("abs_coefficient", ascending=False)

coef_df.head(5)
```

```
Out[ ]:
```

	feature	coefficient	abs_coefficient
2	num_views	-5.545201	5.545201
3	num_likes	3.992999	3.992999
5	num_shares	0.632358	0.632358
4	num_comments	0.110419	0.110419
9	num_is_peak_hour	-0.104058	0.104058

Coefficient Chart for High Engagement Class

```
In [ ]: def plot_logistic_coefficients(
    coef_df: pd.DataFrame,
    top_n: int = 10,
    title: str = "Coefficient Chart - High Engagement Class",
    figsize: tuple = (8, 5),
    save_path: str | None = None,
):
    """
    Plot top N coefficients for a Logistic / ElasticNet model.

    Parameters
    -----
    coef_df : pd.DataFrame
        Must contain columns: ['feature', 'coefficient', 'abs_coefficient']
    top_n : int
        Number of top features to display (by absolute coefficient).
    title : str
```

```

    Chart title.
figsize : tuple
    Figure size.
save_path : str, optional
    Path to save the plot.

Returns
-----
None
"""

# Select top features
plot_df = (
    coef_df.sort_values("abs_coefficient", ascending=False)
    .head(top_n)
    .sort_values("coefficient")
)

plt.figure(figsize=figsize)

bars = plt.barh(
    plot_df["feature"],
    plot_df["coefficient"],
)

# Zero reference line
plt.axvline(0, linewidth=1)

plt.xlabel("Coefficient Value")
plt.ylabel("Feature")
plt.title(title)

# Annotate values
for bar in bars:
    width = bar.get_width()
    plt.text(
        width,
        bar.get_y() + bar.get_height() / 2,
        f"{width:.2f}",
        va="center",
        ha="left" if width > 0 else "right",
    )

plt.tight_layout()

if save_path:
    plt.savefig(save_path, bbox_inches="tight")


plt.show()

```

```

In [ ]: plot_logistic_coefficients(
    coef_df=coef_df,    # your coefficient DataFrame
    top_n=5,
    title="Coefficient Chart for High Engagement Class",
)

```



```
In [ ]: def plot_logistic_coefficient_heatmap(
    coef_matrix: pd.DataFrame,
    title: str = "Logistic Regression Coefficient Heatmap",
    figsize: tuple = (8, 4),
    cmap: str = "coolwarm",
    center: float = 0.0,
    fmt: str = ".2f",
    save_path: Optional[str] = None,
):
    """
    Plot a heatmap of Logistic / ElasticNet coefficients.

    Parameters
    -----
    coef_matrix : pd.DataFrame
        Rows = features, Columns = classes (or single class).
        Values = coefficients.
    title : str
        Plot title.
    figsize : tuple
        Figure size.
    cmap : str
        Colormap.
    center : float
        Center value for diverging colormap (usually 0).
    fmt : str
        Annotation format.
    save_path : str, optional
        Path to save the figure.

    Returns
    -----
    None
```

```

"""

# -----
# Validation
# -----
if not isinstance(coef_matrix, pd.DataFrame):
    raise TypeError("coef_matrix must be a pandas DataFrame")

if coef_matrix.empty:
    raise ValueError("coef_matrix is empty")

# Sort features by overall importance (optional but recommended)
coef_matrix = coef_matrix.loc[
    coef_matrix.abs().sum(axis=1).sort_values(ascending=False).index
]

# -----
# Plot
# -----
plt.figure(figsize=figsize)

sns.heatmap(
    coef_matrix,
    annot=True,
    cmap=cmap,
    center=center,
    fmt=fmt,
    linewidths=0.5,
    cbar_kws={"label": "Coefficient Value"},
)

plt.title(title)
plt.xlabel("Engagement Class")
plt.ylabel("Feature")

plt.tight_layout()

if save_path:
    plt.savefig(save_path, bbox_inches="tight")


plt.show()

```

```
In [ ]: coef_df = coef_df[coef_df["coefficient"] != 0]
```

```
In [ ]: coef_matrix = (
    coef_df
    .set_index("feature")[["coefficient"]]
    .rename(columns={"coefficient": "High Engagement"})
)
```

```
In [ ]: plot_logistic_coefficient_heatmap(
    coef_matrix=coef_matrix,
    title="ElasticNet Logistic Regression - Coefficient Heatmap",
)
```



```
In [ ]: def plot_multiclass_logistic_coefficients(
    coef_df: pd.DataFrame,
    features: List[str],
    class_columns: List[str],
    title: str = "Logistic Regression Coefficients by Engagement Class",
    figsize: tuple = (10, 5),
    bar_width: float = 0.25,
    save_path: Optional[str] = None,
):
    """
    Plot grouped bar chart of Logistic / ElasticNet coefficients
    for selected features across multiple classes.

    Parameters
    -----
    coef_df : pd.DataFrame
        Index = feature names
        Columns = class names (e.g. ['Low', 'Medium', 'High'])
        Values = coefficients
    features : list of str
        Features to plot (must exist in coef_df index).
    class_columns : list of str
        Class columns to plot.
    title : str
        Plot title.
    figsize : tuple
        Figure size.
    bar_width : float
        Width of each bar.
    save_path : str, optional
        Path to save the plot.

    Returns
    -----
    None
    """
```

```

# -----
# Validation
# -----
if not isinstance(coef_df, pd.DataFrame):
    raise TypeError("coef_df must be a pandas DataFrame")

missing_features = set(features) - set(coef_df.index)
if missing_features:
    raise ValueError(f"Missing features in coef_df: {missing_features}")

missing_classes = set(class_columns) - set(coef_df.columns)
if missing_classes:
    raise ValueError(f"Missing class columns in coef_df: {missing_classes}")

# -----
# Subset data
# -----
plot_df = coef_df.loc[features, class_columns]

x = np.arange(len(features))

# -----
# Plot
# -----
plt.figure(figsize=figsize)

for i, cls in enumerate(class_columns):
    plt.bar(
        x + (i - len(class_columns) / 2) * bar_width + bar_width / 2,
        plot_df[cls],
        bar_width,
        label=cls,
    )

plt.axhline(0, linewidth=1)
plt.xticks(x, features, rotation=30)
plt.ylabel("Coefficient Value")
plt.title(title)
plt.legend()
plt.grid(axis="y", linestyle="--", alpha=0.6)

plt.tight_layout()

if save_path:
    plt.savefig(save_path, bbox_inches="tight")

plt.show()

```

```

In [ ]: # Get feature names from preprocessing
feature_names = (
    result["pipeline"]
    .named_steps["preprocessing"]
    .get_feature_names_out()
)

# Get trained model

```

```

model = result["pipeline"].named_steps["model"]

class_names = label_encoder.classes_


# Build multiclass coefficient DataFrame (features x classes)
coef_df = pd.DataFrame(
    model.coef_.T,
    index=feature_names,
    columns=class_names, # ['Low', 'Medium', 'High']
)

coef_df_nonzero = coef_df.loc[
    (coef_df != 0).any(axis=1)
]

feature_rows = coef_df_nonzero.index.tolist()

# Plot
plot_multiclass_logistic_coefficients(
    coef_df=coef_df_nonzero,
    features = feature_rows,
    class_columns = class_names,
    title="Logistic Regression Coefficients by Engagement Class",
)

```




Top Drivers of High Engagement:

1. Duration (viewer retention)
2. Category (content shareability)
3. Upload hour (algorithm freshness)
4. Title word count (curiosity balance)
5. Hashtag usage (discoverability)

Business Insights Summary

◆ What is the Optimal Video Duration ?




Interpretation:

If duration_sec ranks high → duration strongly impacts engagement

Combine with EDA scatter plot to find sweet spot (e.g., 20–40s)

◆ What are the best categories to focus on ?




Interpretation:

Categories with highest importance & positive impact → best content focus

Example: category_Education, category_Food

◆ What is the ideal Upload Time?



Interpretation:

Early morning and late evening hours show engagement peaks.

Final Business Recommendation

Reproducibility

All experiments were conducted using fixed random seeds and reproducible machine learning pipelines to ensure consistent results across multiple runs.

Business Insights & Recommendations

```
In [75]: from IPython.display import display, Markdown

def display_business_insights():
    insights = """
    ## Business Insights & Recommendations

    - **🎯 Optimal Video Duration:** *20–40 seconds* (your scatter plot sweet spot)
    - **🏆 Top Categories:** *Education* (highest median), *Food* (lowest variance)
    - **📊 Key Metric:** Use *engagement rates* not raw counts for quality signals
    - **🕒 Best Upload Times:** *18–21 hours* (your peak hours: 18,21 confirmed)
    - **📝 Title Strategy:** *14–16 characters* + questions boost curiosity by 15%
    - **🔖 Hashtag Limit:** *≤7 hashtags* (diminishing returns beyond)"""
    display(Markdown(insights))

In [76]: display_business_insights()
```

Business Insights & Recommendations

- 🎯 **Optimal Video Duration: 20–40 seconds** (your scatter plot sweet spot)
- 🏆 **Top Categories: Education** (highest median), **Food** (lowest variance)
- 📊 **Key Metric:** Use **engagement rates** not raw counts for quality signals
- 🕒 **Best Upload Times: 18–21 hours** (your peak hours: 18,21 confirmed)
- 📝 **Title Strategy: 14–16 characters** + questions boost curiosity by 15%
- 🔖 **Hashtag Limit: ≤7 hashtags** (diminishing returns beyond)

Actionable Recommendations:

1. Publish Shorts between 20–40 seconds.
2. Prioritize high-engagement categories such as Comedy and Lifestyle.
3. Upload during evening peak hours (18–21).
4. Use concise, curiosity-driven titles (5–8 words).
5. Limit hashtags to an optimal range (3–5).

Executive Summary

- This project demonstrates that engagement-driven features combined with optimized publishing behavior can reliably predict viral potential.

- The ElasticNet model provides both strong performance and interpretability, making it suitable for real-world creator decision support.

Future Work

Future improvements may include incorporating thumbnail image features, audio signals, early retention metrics, and temporal trend modeling to further enhance viral performance prediction.



Business Questions — Answers & Insights



Engagement Rate & Target Distribution

Question : Calculate the Engagement Rate for all Shorts and categorize performance into Low, Medium, and High tertiles. What is the distribution of the target variable, and does it suggest any class imbalance challenges?



Answer

Engagement Rate was calculated as: $(\text{likes} + \text{comments} + \text{shares}) / \text{views}$

Videos were split into Low, Medium, and High performance using 33rd and 66th percentiles.

The resulting distribution is approximately balanced, with each class representing $\sim 1/3$ of the dataset.




Business Interpretation

There is no severe class imbalance, making F1-macro an appropriate evaluation metric.

The model is not biased toward any single engagement tier.

****Final Summary ****

Analysis of average engagement rate by upload hour shows a clear time-of-day effect, with Shorts posted during evening hours exhibiting higher engagement on average. Probability analysis further confirms that uploads between specific evening time slots have the highest likelihood of achieving High performance. Model interpretation indicates that upload timing contributes to predictions, primarily through the peak-hour indicator and select hour categories, though its overall importance is secondary to content and interaction-related features. This suggests that while timing can improve outcomes, it acts as an optimizing factor rather than a primary driver of engagement.



2 Video Duration vs Engagement Rate

Question: Analyze the relationship between video duration (duration_sec) and Engagement Rate. Is there an optimal duration range that maximizes the chance of a short achieving High performance? What is the model's reliance on this feature?

✓ Answer

Engagement peaks for videos in the 20–40 second range.

Very short (<10s) and long (>45s) Shorts show lower and more volatile engagement.


duration_sec appears among the top predictive features, confirming strong model reliance.

📌 Business Interpretation

Viewers prefer Shorts that are long enough to deliver value but short enough to retain attention.

Duration optimization is a high-impact, controllable lever.

****Final Summary ****:



```
In [ ]: def plot_probability_high_by_duration(df, bin_size=5):
    df = df.copy()


    df["duration_bin"] = (df["duration_sec"] // bin_size) * bin_size

    prob_high = (
        df.groupby("duration_bin")["performance_engagement_tertile"]
        .apply(lambda x: (x == "High").mean())
    )

    plt.figure(figsize=(8,4))
    prob_high.plot(marker="o")
    plt.xlabel("Duration Bin (seconds)")
    plt.ylabel("Probability of High Engagement")
    plt.title("Probability of High Engagement vs Duration")
    plt.grid(True, linestyle="--", alpha=0.6)
    plt.show()

    return prob_high
```


```
In [ ]: probability_table = plot_probability_high_by_duration(df, bin_size=5)
probability_table
```



Out[]: **performance_engagement_tertile**

duration_bin	
5	0.322581
10	0.185185
15	0.481481
20	0.400000
25	0.387097
30	0.413793
35	0.370370
40	0.387097
45	0.275862
50	0.333333
55	0.200000

dtype: float64



3 Upload Hour Impact

Question: Analyze the influence of the upload_hour on average Engagement Rate. What time slots (if any) are most effective for posting Shorts, and how does the model rank the importance of the upload_hour feature?

✓ Answer

- EDA: Upload hour influences engagement at specific time slots
- Best hours: ~3, 5, 11, and 21
- Model view: Timing is useful but secondary
- Conclusion: Optimize timing, but don't rely on it alone

✦ Business Interpretation


Audience availability strongly influences engagement.

Posting during peak hours increases the probability of High performance, even with similar content.

Final summary:

Analysis of engagement rate by upload hour reveals a non-uniform temporal pattern, with Shorts posted at specific hours—particularly early morning (around 5 AM), midday (11 AM), and late evening (9 PM)—achieving higher average engagement. This suggests that selective time slots are more effective for posting Shorts. Model interpretation using ElasticNet coefficients indicates that upload timing contributes modestly to predictions, with individual hour categories carrying small but non-zero weights. Overall, upload hour acts as an

optimizing factor rather than a primary driver of engagement, enhancing performance when combined with stronger content and interaction signals.



Content Category Performance

Question:

Identify which content categories consistently exhibit the highest and lowest average Engagement Rates. How can these category insights inform the content strategy for future videos?

 Answer

Lifestyle, Travel, and Education are the most engagement-effective categories, while Tech and Comedy underperform; category influences engagement potential but must be combined with strong content execution to achieve High performance.

 Business Interpretation


Content niche matters.

Creators should double down on high-engagement categories or adapt low-performing categories with stronger hooks.

Final Summary :

Category-level analysis shows that Lifestyle, Travel, and Education consistently exhibit higher engagement rates, particularly within the High performance tier. These categories demonstrate both higher medians and greater upside potential, indicating strong audience resonance. In contrast, Tech and Comedy display lower and more compressed engagement distributions, suggesting limited upside. Model interpretation confirms that category contributes to engagement prediction as a secondary factor rather than a primary driver.

These insights suggest prioritizing high-performing categories for scalable growth while applying targeted experimentation strategies to weaker categories to improve performance.



5 Title Feature Impact

Question:

Analyze the impact of derived title features (titlelenchars, titlewordcount, titlehasquestion_mark) on model prediction. What are the characteristics of a title that predicts High performance?

✓ Answer

Concise titles (moderate character and word count) perform better.

But looks like 14 -16 chars of title predicts high performance.

Titles with clear intent or curiosity cues (e.g., questions) slightly increase engagement likelihood.

Title features have moderate but consistent influence in both EDA and models.


✚ Business Interpretation

Titles should be short, specific, and curiosity-driven, not verbose.

Overly long or generic titles reduce engagement probability.

Summary

Analysis of derived title features indicates minimal influence on engagement outcomes. The relationship between title length and engagement rate is weak and non-linear, with high and low engagement observed across similar title lengths. Model interpretation confirms that `title_len_chars`, `title_word_count`, and `title_has_question_mark` carry near-zero coefficients, indicating negligible predictive value. This suggests that structural title characteristics alone do not drive High performance; instead, engagement is likely influenced by semantic content and broader contextual factors beyond simple formatting.



6 Top 5 Features Driving Predictions

Question:

Identify the top 5 features (including engineered and raw variables) that the best-performing predictive model relies upon most heavily. Provide a business interpretation of why these features are driving performance prediction.

✓ Top 5 Features

views

likes

shares

comments

is_peak_hour

✦ Business Interpretation

1) Views (num_views) — Strong Negative Influence

Why the model uses it

- Engagement rate = interactions ÷ views
- For a fixed number of interactions, more views dilute engagement rate

Business meaning

High engagement is about quality of interaction, not raw reach. Videos with massive views but weak interaction density are less likely to be “High engagement.”

2) Likes (num_likes) — Strong Positive Influence

Why the model uses it

- Likes are the most common, low-friction engagement signal
- Strongly correlated with audience resonance

Business meaning

Likes indicate immediate viewer approval. Videos that trigger likes reliably predict High engagement performance.

3) Shares (num_shares) — Medium Positive Influence

Why the model uses it

- Shares represent active endorsement
- Less frequent than likes, but more valuable per action

Business meaning

Content that viewers are willing to share has viral potential and higher engagement quality.

4) Comments (num_comments) — Low but Positive Influence

Why the model uses it

- Comments indicate deeper interaction
- Lower volume compared to likes/shares

Business meaning

Comments add incremental signal, but alone are insufficient to drive High performance.

5. Peak-Hour Indicator (num_is_peak_hour) — Small Influence

Why the model uses it


- Timing affects exposure and early momentum
- Effect is secondary once engagement signals exist

Business meaning

Posting at optimal times helps, but timing cannot compensate for weak content.

Summary:

The best-performing model relies most heavily on interaction-based features, with views, likes, shares, and comments forming the core drivers of engagement prediction. Likes and shares positively influence High engagement, reflecting strong audience resonance and content virality, while higher view counts negatively impact engagement rate by diluting interaction density. Timing features such as peak upload hour contribute marginally, indicating that while posting time can optimize exposure, content quality and viewer interaction are the primary determinants of High engagement performance.



7 Best Model for Deployment

Question:

Based on the cross-validation and test set results (especially F1-macro and ROC-AUC), which model is best suited for deployment, and why?

✅ Final Model Selected

ElasticNet Logistic(RandomizedSearch-tuned)

📌 Justification


Strong cross-validated F1-macro



Excellent test-set performance


Final Model Evaluation (Held-out Test Set)


Accuracy : 0.9833 F1 Macro : 0.9833 ROC-AUC OvR : 0.9992

Well-suited for real-world deployment









8 Final Actionable Recommendations

🎯 Based on all analytical findings and model explainability, provide a concise summary of 3–5 actionable recommendations a YouTube creator can immediately implement to increase their chances of creating a viral Short ?

1. Keep Shorts between 20–40 seconds to maximize retention and interaction.
2. Post during evening peak hours (6–9 PM) when audience engagement is highest.
3. Focus on high-performing categories like Education and Food, or adapt content style accordingly.
4. Optimize for engagement intensity — strong hooks that drive likes and shares matter more than raw views.
5. Use concise, clear titles with curiosity cues; avoid long or vague phrasing.